

CS202 : COMPUTER ORGANIZATION

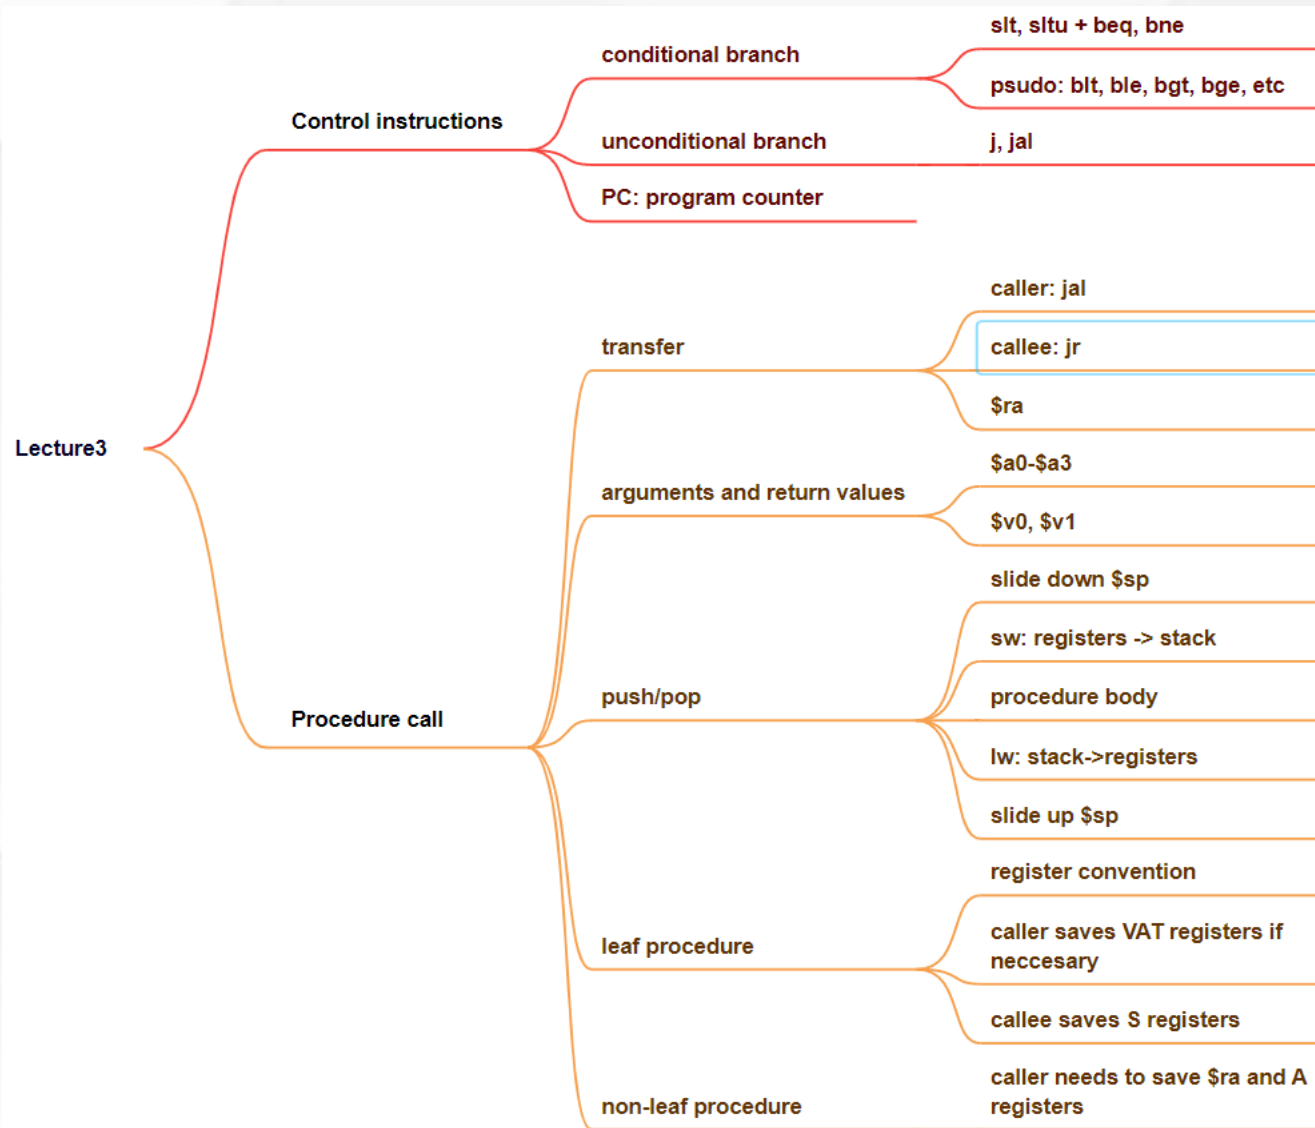
Lecture 4 Instruction Set Architecture (3)

2023 Spring

Today's Agenda

- Recap
 - More control instructions
 - Procedure Call
- Context
 - MIPS addressing
 - Translating and starting a program
 - Other popular ISAs (RISC-V)
- Reading: Textbook 2.10, 2.12-13, 2.19
- Self-study: The rest sections of Chapter 2


Recap



Memory Layout summary

- Text: program code
- Static data: global variables
 - e.g., static variables in C, constant arrays and strings
 - \$gp initialized to address allowing \pm offsets into this segment
- Dynamic data: heap
 - E.g., malloc in C, new in Java
- Stack: automatic storage

High address



Low address

MIPS Addressing


- Addressing: how the instructions identify the operands of the instruction.
- MIPS Addressing mode:
 - Register addressing
 - Immediate addressing
 - Base/Displacement addressing
 - PC-relative addressing
 - Pseudo-direct addressing

R-format → Register Addressing

- Using register as the operand
- E.g.,
 - Arithmetic: `add`, `sub`...
 - Logical: `and`, `or`, `sll`, `srl`...
 - Control: `slt`...



I-format → Immediate Addressing



- For instructions including immediate
 - E.g. `addi`, `andi`, `ori`
- Most constants are small
 - 16-bit immediate is sufficient
- For the occasional 32-bit constant
 - `lui rt, constant` (Load upper immediate)
 - Copies 16-bit constant to left 16 bits of `rt`
 - Clears right 16 bits of `rt` to 0

`lui $s0, 61`

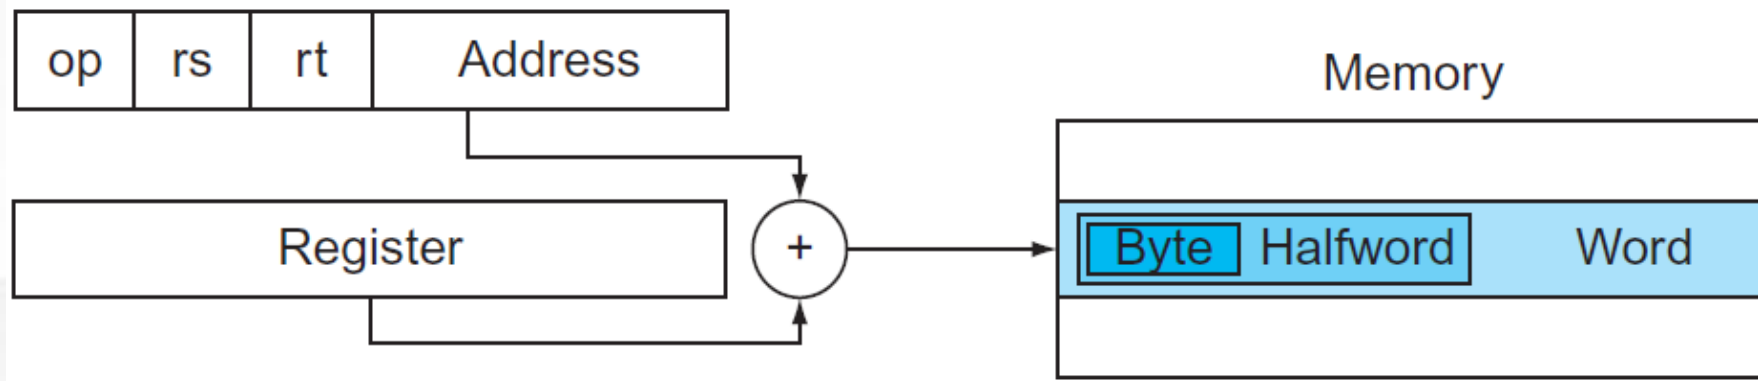
0000 0000 0011 1101	0000 0000 0000 0000
---------------------	---------------------

`ori $s0, $s0, 2304`

0000 0000 0011 1101	0000 1001 0000 0000
---------------------	---------------------

I-format → Base Addressing

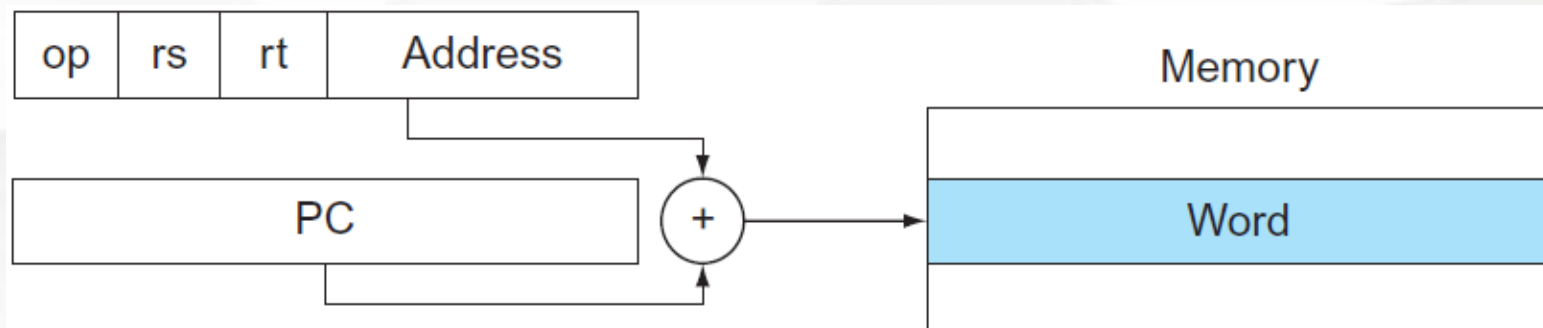
- Also called Displacement Addressing
- The operand is at the memory, whose address is the sum of a register and a constant **lw/lh/lb/sw/sh/sb**
 - e.g. **lw \$s0, 4(\$s1)**
 - op: 100011, rs: 10001 (address of s1), rt: 10000 (address of s0), offset: 100 (4)
 - Instruction:



I-format → PC-relative addressing(1)

- Branch instructions specify
 - Opcode, two registers, target address, `beq/bne`
 - e.g. `beq $s0, $s1, label`
- Most branch targets are near branch
 - Forward or backward
- PC-relative addressing
 - Target address = $PC + \text{offset} \times 4$
 - PC already incremented by 4 by this time


Addr.	Inst.
0	beq \$t0,\$t1,L1
4	lw ...
8	add...
12	L1:...



I-format → PC-relative addressing(2)

- Branch Addressing: Use the address field as a two's complement word offset to PC
 - Branches generally change the PC by a small amount
 - Can specify $\pm 2^{15}$ aligned word addresses from the PC
- Why not use byte address offset from PC?
 - MIPS Instructions are “word-aligned”: Address is always a multiple of 4 (in bytes)
 - PC always points to an instruction
 - Instead of specifying $\pm 2^{15}$ bytes from the PC, we will now specify $\pm 2^{15}$ words = $\pm 2^{17}$ byte addresses around PC

J-format → Pseudodirect addressing(1)




- Jump (**j** and **jal**) targets could be anywhere in text segment
 - Encode full address in instruction
- Addressing in jumps
 - why absolute addressing fails?
 - solution #1: count in words
 - solution #2: invent a new format (J-type)
 - solution #3: borrow 4-bits from **PC**



J-format → Pseudodirect addressing(2)

- Jump Addressing

- the jump address is the 26 bits of the instruction concatenated with the upper bits of the PC
- Target address = $PC[31...28] : (\text{address} \times 4)$



- limit on code size (address boundary)

- 26 bits word address = 28 bits byte address
= 64 MWords = 256 MB

Target Addressing Example

- C code

```
while(save[i] == k)
    i += 1;
```


Register assignment:

i(\$s3), k (\$s5), address of save (\$s6)

- MIPS code

Loop: sll \$t1, \$s3, 2	# Temp reg \$t1 = i*4
add \$t1, \$t1, \$s6	# \$t1 = address of save[i]
lw \$t0, 0(\$t1)	# Temp reg \$t0 = save[i]
bne \$t0, \$s5, Exit	# go to Exit if save[i]!=k
addi \$s3, \$s3, 1	# i=i+1
j Loop	# go to Loop

Exit: ...



Target Addressing Example

- **bne**

op	rs	rt	address
6 bits	5 bits	5 bits	16 bits

bne instr. target addr. = PC+ address x 4
- **j**

op	address
6 bits	26 bits

j instr. target addr. = PC[31...28] :(address x 4)

Assume Loop at location 80000

Loop: sll	\$t1, \$s3, 2	80000	0	0	19	9	4	0
add	\$t1, \$t1, \$s6	80004	0	9	22	9	0	32
lw	\$t0, 0(\$t1)	80008	35	9	8	0		
bne	\$t0, \$s5, Exit	80012	5	8	21	?		
addi	\$s3, \$s3, 1	80016	8	19	19	1		
j	Loop	80020	2	?				
Exit: ...		80024						

- $80024 = PC+4 + \text{address} \times 4 = 80012 + 4 + 2 \times 4$
- $80000 = PC[31...28] : (\text{address} \times 4) = 80020[31...28] : (\text{address} \times 4) = 0 : 20000 \times 4$

Branching Far Away

- If branch target is too far to encode with 16-bit offset, assembler rewrites the code
- Example

```
        beq $s0,$s1, L1
          ↓
        bne $s0,$s1, L2
        j   L1
L2:      ...
```


1. Immediate addressing




2. Register addressing




3. Base addressing



4. PC-relative addressing



5. Pseudodirect addressing



Decoding Machine Code

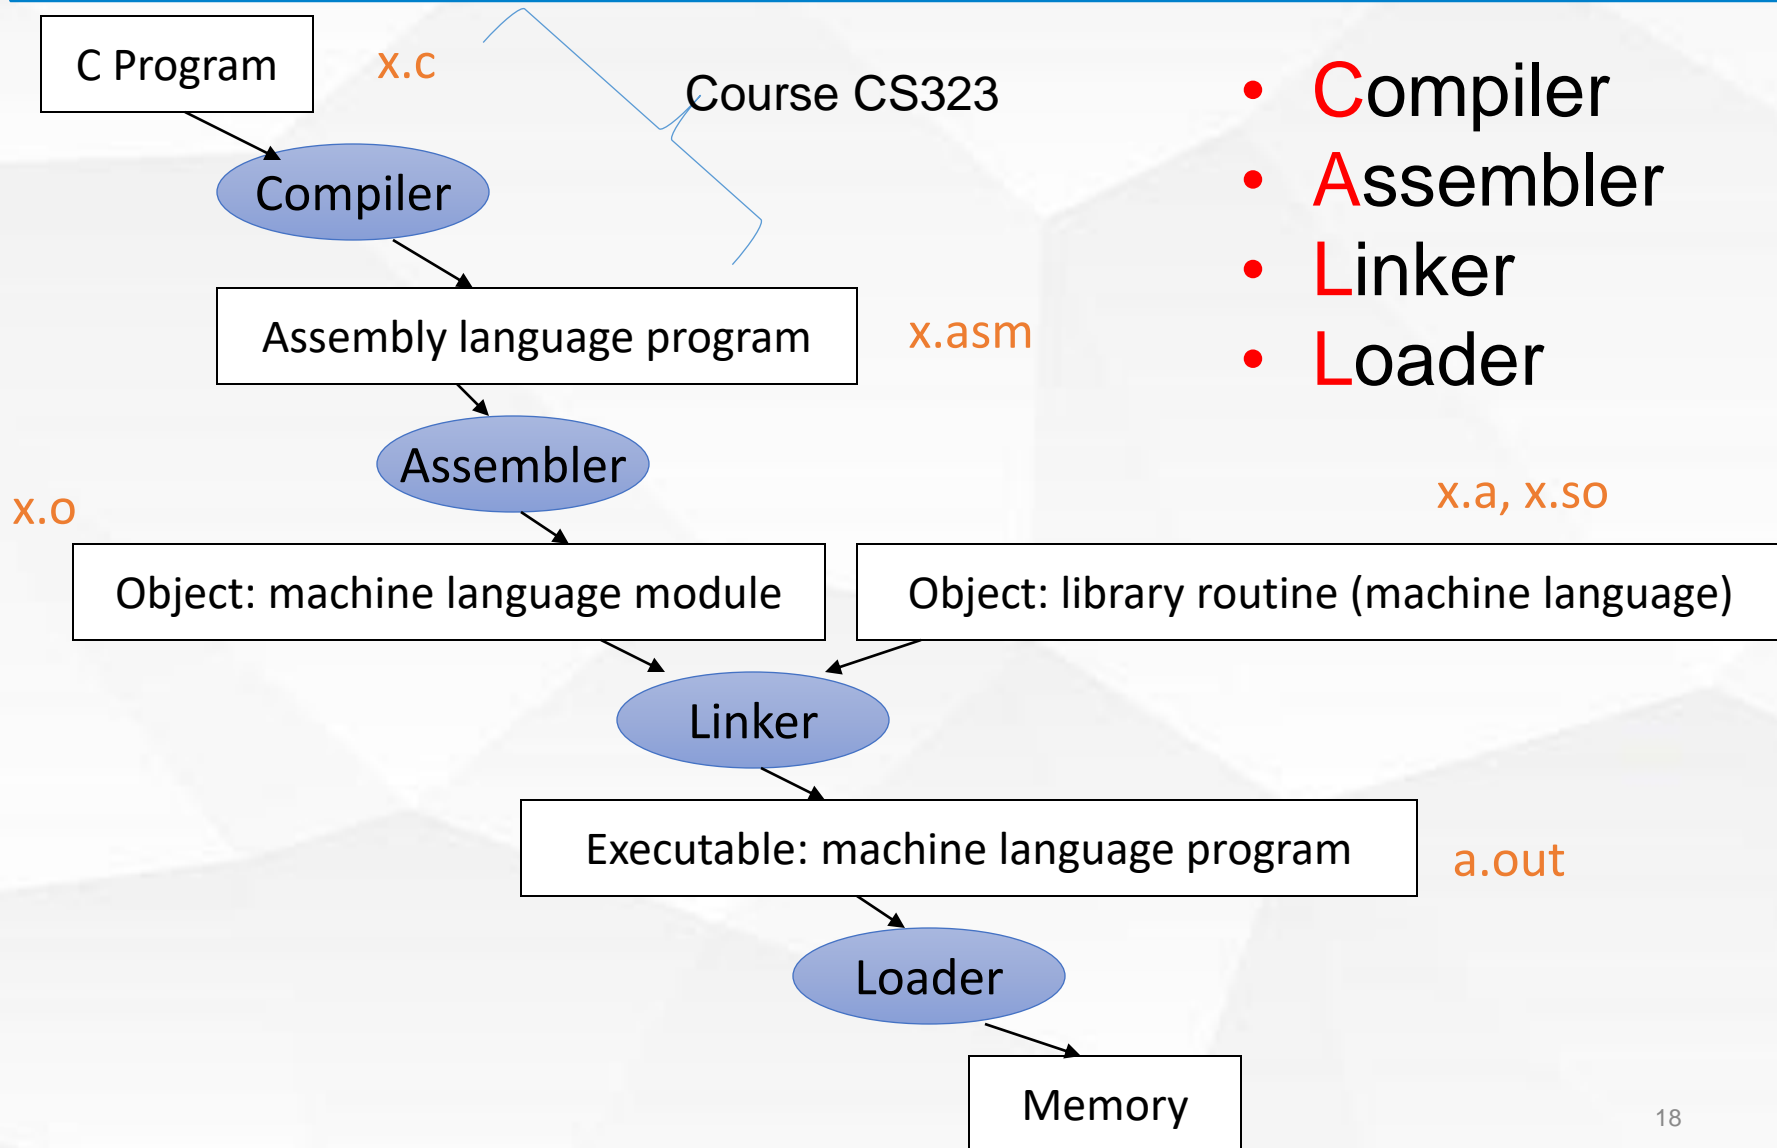
- What is the assembly language of the following machine instruction?

0x00af8020

- Hex to bin: 0000 0000 1010 1111 1000 0000 0010 0000
 - op 000000 → R-format
 - rs 00101 → 5₁₀
 - rt 01111 → 15₁₀
 - rd 10000 → 16₁₀
 - shamt 00000
 - func 100000
- Get the instruction: add \$s0, \$a1, \$t7

Name	Fields						Comments
Field size	6 bits	5 bits	5 bits	5 bits	5 bits	6 bits	All MIPS instructions are 32 bits long
R-format	op	rs	rt	rd	shamt	funct	Arithmetic instruction format
I-format	op	rs	rt	address/immediate			Transfer, branch, imm. format
J-format	op	target address					Jump instruction format 17

Running a Program: *CALL*



Translation vs. Interpretation

- How do we run a program written in a source language?
 - Interpreter: Directly executes a program in the source language, e.g. Python
 - closer to high-level, but slower
 - Translator: Converts a program from the source language to an equivalent program in another language, e.g. C
 - more efficient
- Translation/compilation helps “hide” the program “source” from the users
 - One model for creating value in the marketplace (e.g. Microsoft keeps all their source code secret)
 - Alternative model, “open source”, creates value by publishing the source code and fostering a community of developers

Compiler

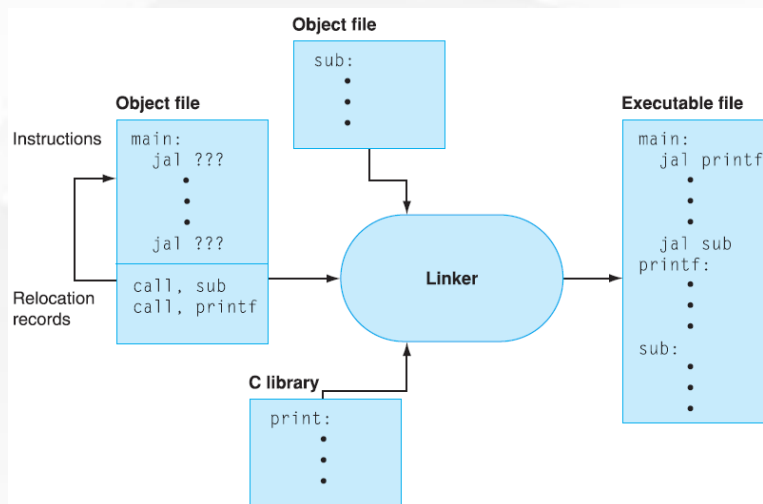
- Input: Higher-level language code
 - e.g. C, Java in files such as foo.c
- Output: Assembly Language Code
 - e.g. foo.asm
- The output may contain pseudo-instructions
- Optimizing compilers today can produce assembly language programs nearly as well as an assembly language expert, and sometimes even better for large programs.

Assembler

- Input: Assembly language code
 - e.g. foo.asm
- Output: Object code (True Assembly), information tables
 - e.g. foo.o – Object file
- Convert pseudo-instructions into actual hardware instructions
 - e.g. `move $t0, $t1` → `add $t0, $zero, $t1`
`blt $t0, $t1, L` → `slt $at, $t0, $t1`
`bne $at, $zero, L`
`$at` (register1) reserved for assembler
- Convert assembly instructions into machine instructions
 - compute the actual values for instruction labels
 - maintain info on external references and debugging information

Linker

- Input: Object Code files, information tables
 - e.g. foo.o, lib.o
- Output: Executable Code
 - e.g. a.out
- Stitches different object files into a single executable
 - patch internal and external references
 - determine addresses of data and instruction labels
 - organize code and data modules in memory



Loader

- Input: Executable Code
 - e.g. a.out
- Output: <program is run>
- Executable files are stored on disk
- When one is run, loader's job is to load it into memory and start it running
- In reality, loader is the operating system (OS)
 - loading is one of the OS tasks

Full Example: Sort in C

- Allocate registers to program variables
- Produce code for the program body
- Preserve registers across procedure invocations

```
void sort (int v[ ], int n)
{
    int i, j;
    for (i=0; i<n; i+=1) {
        for (j=i-1; j>=0 && v[j] > v[j+1]; j-=1) {
            swap (v,j);
        }
    }
}
```

```
void swap (int v[ ], int k)
{
    int temp;
    temp = v[k];
    v[k] = v[k+1];
    v[k+1] = temp;
}
```


The swap Procedure

- Register allocation: \$a0 and \$a1 for the two arguments, \$t0 for the temp variable – no need for saves and restores as we're not using \$s0-\$s7 and this is a leaf procedure (won't need to re-use \$a0 and \$a1)

```
swap:  sll  $t1, $a1, 2
       add  $t1, $a0, $t1
       lw   $t0, 0($t1)
       lw   $t2, 4($t1)
       sw   $t2, 0($t1)
       sw   $t0, 4($t1)
       jr   $ra
```

```
void swap (int v[], int k)
{
    int temp;
    temp = v[k];
    v[k] = v[k+1];
    v[k+1] = temp;
}
```

The Procedure Body

```
void sort (int v[ ], int n) {
    int i, j;
    for (i=0; i<n; i+=1) {
        for (j=i-1; j>=0 && v[j] > v[j+1]; j-=1) {
            swap (v,j);
        }
    }
}
```





	move \$s2, \$a0	# save \$a0 into \$s2	Move params
	move \$s3, \$a1	# save \$a1 into \$s3	
for1tst:	move \$s0, \$zero	# i = 0	Outer loop
	slt \$t0, \$s0, \$s3	# \$t0 = 0 if \$s0 ≥ \$s3 (i ≥ n)	
	beq \$t0, \$zero, exit1	# go to exit1 if \$s0 ≥ \$s3 (i ≥ n)	
	addi \$s1, \$s0, -1	# j = i - 1	
for2tst:	slti \$t0, \$s1, 0	# \$t0 = 1 if \$s1 < 0 (j < 0)	
	bne \$t0, \$zero, exit2	# go to exit2 if \$s1 < 0 (j < 0)	
	sll \$t1, \$s1, 2	# \$t1 = j * 4	Inner loop
	add \$t2, \$s2, \$t1	# \$t2 = v + (j * 4)	
	lw \$t3, 0(\$t2)	# \$t3 = v[j]	
	lw \$t4, 4(\$t2)	# \$t4 = v[j + 1]	
	slt \$t0, \$t4, \$t3	# \$t0 = 0 if \$t4 ≥ \$t3	
	beq \$t0, \$zero, exit2	# go to exit2 if \$t4 ≥ \$t3	
	move \$a0, \$s2	# 1st param of swap is v (old \$a0)	Pass params & call
	move \$a1, \$s1	# 2nd param of swap is j	
	jal swap	# call swap procedure	
	addi \$s1, \$s1, -1	# j -= 1	Inner loop
	j for2tst	# jump to test of inner loop	
exit2:	addi \$s0, \$s0, 1	# i += 1	
	j for1tst	# jump to test of outer loop	Outer loop

The Full Procedure




sort:	addi \$sp,\$sp, -20	# make room for 5 registers
	sw \$ra, 16(\$sp)	# save \$ra on stack
	sw \$s3,12(\$sp)	# save \$s3 on stack
	sw \$s2, 8(\$sp)	# save \$s2 on stack
	sw \$s1, 4(\$sp)	# save \$s1 on stack
	sw \$s0, 0(\$sp)	# save \$s0 on stack
	...	# procedure body
	...	
	exit1: lw \$s0, 0(\$sp)	# restore \$s0 from stack
	lw \$s1, 4(\$sp)	# restore \$s1 from stack
	lw \$s2, 8(\$sp)	# restore \$s2 from stack
	lw \$s3,12(\$sp)	# restore \$s3 from stack
	lw \$ra,16(\$sp)	# restore \$ra from stack
	addi \$sp,\$sp, 20	# restore stack pointer
	jr \$ra	# return to calling routine

Effect of Compiler Optimization

- Compiled with gcc for Pentium 4 under Linux



Effect of Language and Algorithm



Lessons Learnt

- Instruction count and CPI are not good performance indicators in isolation
- Compiler optimizations are sensitive to the algorithm
- Java/JIT compiled code is significantly faster than JVM interpreted
 - Comparable to optimized C in some cases
- Nothing can fix a dumb algorithm!

Mainstream ISAs

Full RISC-V Architecture:

<https://digitalassets.lib.berkeley.edu/techreports/ucb/text/EECS-2016-1.pdf>



x86

Designer	Intel, AMD
Bits	16-bit, 32-bit and 64-bit
Introduced	1978 (16-bit), 1985 (32-bit), 2003 (64-bit)
Design	CISC
Type	Register-memory
Encoding	Variable (1 to 15 bytes)
Endianness	Little

Core i3, i5, i7...



ARM architectures

Designer	ARM Holdings
Bits	32-bit, 64-bit
Introduced	1985; 31 years ago
Design	RISC
Type	Register-Register
Encoding	AArch64/A64 and AArch32/A32 use 32-bit instructions, T32 (Thumb-2) uses mixed 16- and 32-bit instructions. ARMv7 <i>user-space</i> compatibility ^[1]
Endianness	Bi (little as default)

Smartphone-like devices
(iPhone, Android), Raspberry Pi, Embedded systems
Apple M series

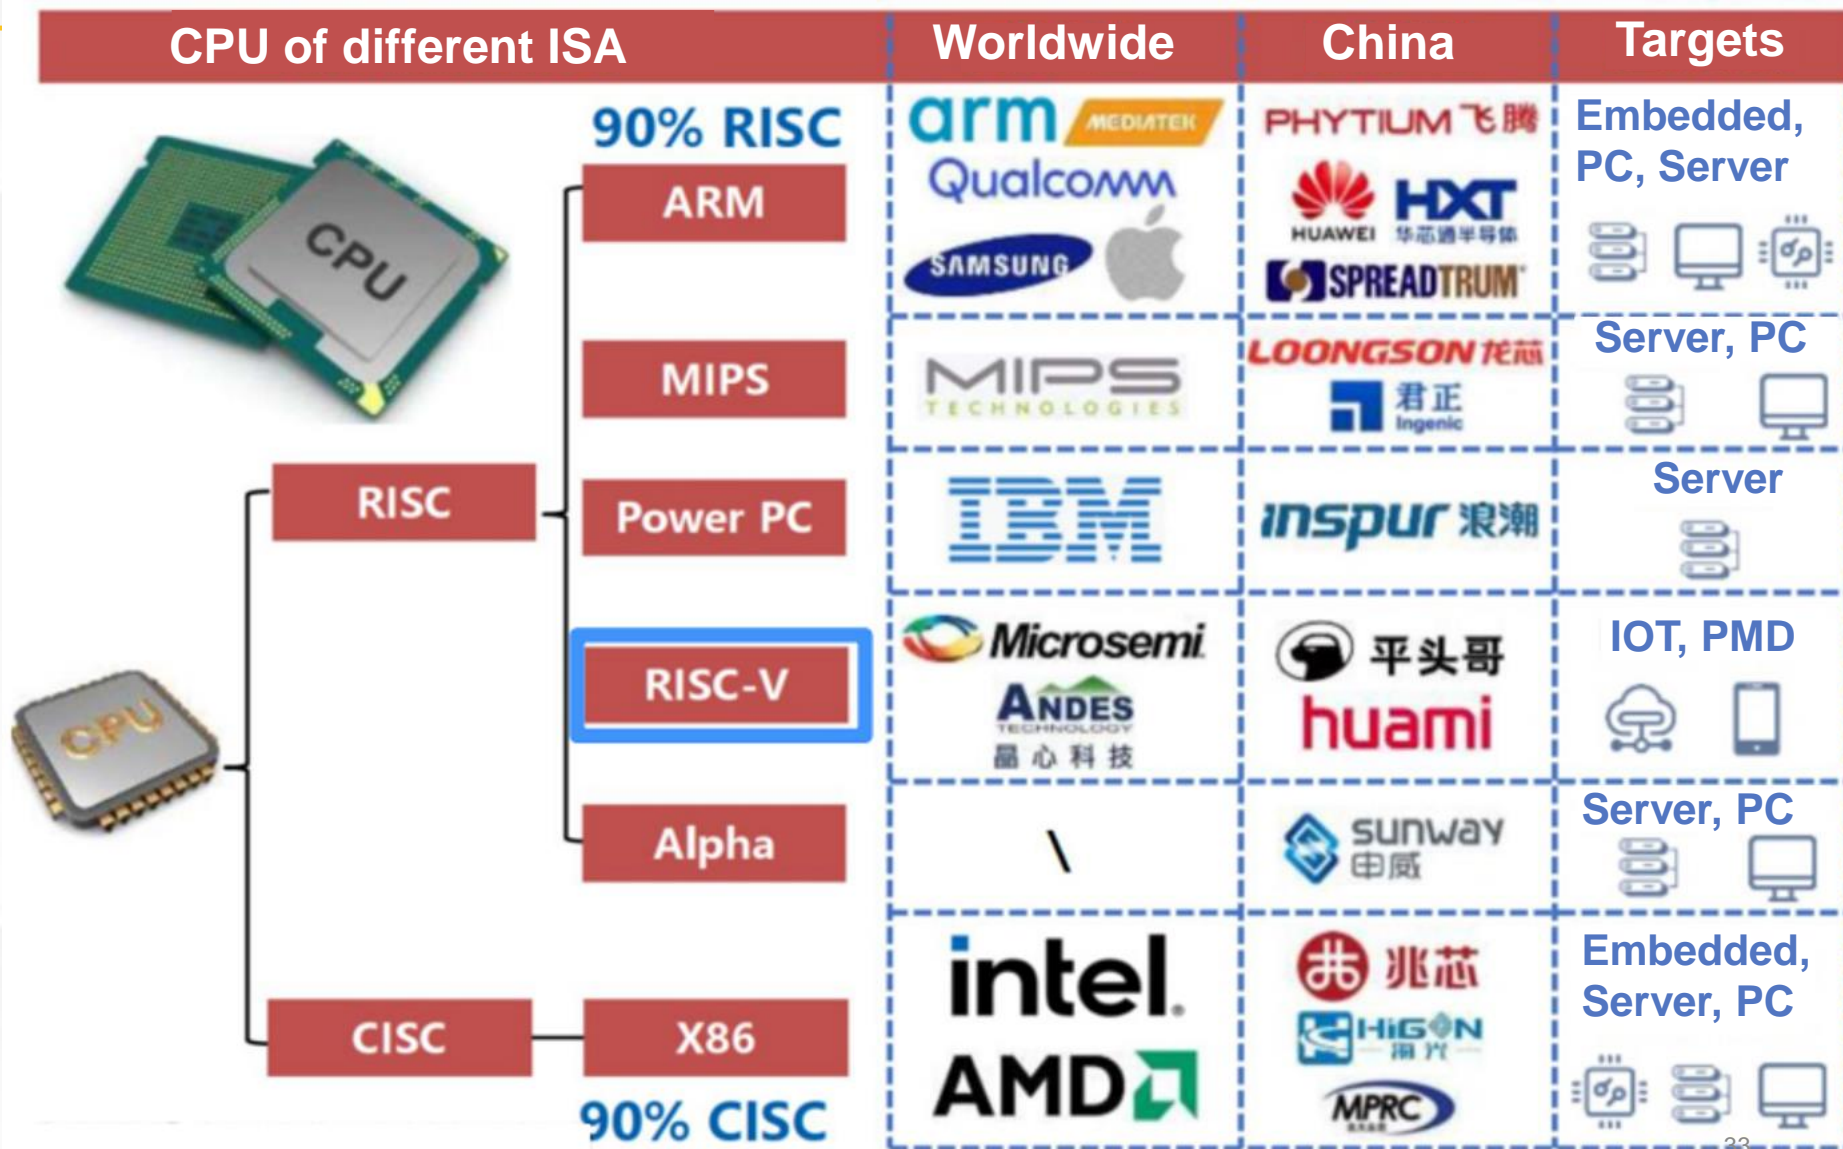


RISC-V

Designer	University of California, Berkeley
Bits	32, 64, 128
Introduced	2010
Version	2.2
Design	RISC
Type	Load-store
Encoding	Variable
Branching	Compare-and-branch
Endianness	Little

Versatile and open-source
Relatively new, designed for cloud computing, embedded systems, academic use


RISC-V among ISAs



Rapid RISC-V growth led by industrial

- Semico Research predicts the market will consume 62.4 billion RISC-V CPU cores by 2025, a 146.2% CAGR 2018-2025. The industrial sector to lead with 16.7 billion cores.
- Custom ICs Based on RISC-V Will Enable Cost-Effective IoT Product Differentiation

--- Gartner, June 2020



RISC-V foundation



Foundation: 100+ Members



MIPS vs. RISC-V

- Similar basic set of instructions

	MIPS32	RISC-V (RV32)
Date announced	1985	2010
License	Proprietary	Open-Source
Instruction size	32 bits	32 bits
Endianness	Big-endian	Little-endian
Addressing modes	5	4
Registers	32 × 32-bit	32 × 32-bit
Pipeline Stages	5 stages	5 stages
ISA type	Load-store	Load-store
Conditional branches	slt, sltu + beq, bnq	+blt,bge,bltu,bgeu

RISC-V Registers

- x0: the constant value 0
- x1: return address
- x2: stack pointer
- x3: global pointer
- x4: thread pointer
- x5 – x7, x28 – x31: temporaries
- x8: frame pointer
- x9, x18 – x27: saved registers
- x10 – x11: function arguments/results
- x12 – x17: function arguments

Name	Register number	Usage	Preserved on call?
x0	0	The constant value 0	n.a.
x1 (ra)	1	Return address (link register)	yes
x2 (sp)	2	Stack pointer	yes
x3 (gp)	3	Global pointer	yes
x4 (tp)	4	Thread pointer	yes
x5-x7	5-7	Temporaries	no
x8-x9	8-9	Saved	yes
x10-x17	10-17	Arguments/results	no
x18-x27	18-27	Saved	yes
x28-x31	28-31	Temporaries	no

RISC-V Assembly

RISC-V assembly language				
Category	Instruction	Example	Meaning	Comments
Arithmetic	Add	add x5, x6, x7	$x5 = x6 + x7$	Three register operands
	Subtract	sub x5, x6, x7	$x5 = x6 - x7$	Three register operands
	Add immediate	addi x5, x6, 20	$x5 = x6 + 20$	Used to add constants
Data transfer	Load doubleword	ld x5, 40(x6)	$x5 = \text{Memory}[x6 + 40]$	Doubleword from memory to register
	Store doubleword	sd x5, 40(x6)	$\text{Memory}[x6 + 40] = x5$	Doubleword from register to memory
	Load word	lw x5, 40(x6)	$x5 = \text{Memory}[x6 + 40]$	Word from memory to register
	Load word, unsigned	lwu x5, 40(x6)	$x5 = \text{Memory}[x6 + 40]$	Unsigned word from memory to register
	Store word	sw x5, 40(x6)	$\text{Memory}[x6 + 40] = x5$	Word from register to memory
	Load halfword	lh x5, 40(x6)	$x5 = \text{Memory}[x6 + 40]$	Halfword from memory to register
	Load halfword, unsigned	lhu x5, 40(x6)	$x5 = \text{Memory}[x6 + 40]$	Unsigned halfword from memory to register
	Store halfword	sh x5, 40(x6)	$\text{Memory}[x6 + 40] = x5$	Halfword from register to memory
	Load byte	lb x5, 40(x6)	$x5 = \text{Memory}[x6 + 40]$	Byte from memory to register
	Load byte, unsigned	lbu x5, 40(x6)	$x5 = \text{Memory}[x6 + 40]$	Byte halfword from memory to register
	Store byte	sb x5, 40(x6)	$\text{Memory}[x6 + 40] = x5$	Byte from register to memory
	Load reserved	lr.d x5, (x6)	$x5 = \text{Memory}[x6]$	Load; 1st half of atomic swap
	Store conditional	sc.d x7, x5, (x6)	$\text{Memory}[x6] = x5; x7 = 0/1$	Store; 2nd half of atomic swap
	Load upper immediate	lui x5, 0x12345	$x5 = 0x12345000$	Loads 20-bit constant shifted left 12 bits
Logical	And	and x5, x6, x7	$x5 = x6 \& x7$	Three reg. operands; bit-by-bit AND
	Inclusive or	or x5, x6, x8	$x5 = x6 x8$	Three reg. operands; bit-by-bit OR
	Exclusive or	xor x5, x6, x9	$x5 = x6 \wedge x9$	Three reg. operands; bit-by-bit XOR
	And immediate	andi x5, x6, 20	$x5 = x6 \& 20$	Bit-by-bit AND reg. with constant
	Inclusive or immediate	ori x5, x6, 20	$x5 = x6 20$	Bit-by-bit OR reg. with constant
	Exclusive or immediate	xori x5, x6, 20	$x5 = x6 \wedge 20$	Bit-by-bit XOR reg. with constant

RISC-V Assembly


RISC-V assembly language				
Category	Instruction	Example	Meaning	Comments
Shift	Shift left logical	<code>sll x5, x6, x7</code>	$x5 = x6 \ll x7$	Shift left by register
	Shift right logical	<code>srl x5, x6, x7</code>	$x5 = x6 \gg x7$	Shift right by register
	Shift right arithmetic	<code>sra x5, x6, x7</code>	$x5 = x6 \ggg x7$	Arithmetic shift right by register
	Shift left logical immediate	<code>slli x5, x6, 3</code>	$x5 = x6 \ll 3$	Shift left by immediate
	Shift right logical immediate	<code>srl i x5, x6, 3</code>	$x5 = x6 \gg 3$	Shift right by immediate
	Shift right arithmetic immediate	<code>srai x5, x6, 3</code>	$x5 = x6 \ggg 3$	Arithmetic shift right by immediate
Conditional branch	Branch if equal	<code>beq x5, x6, 100</code>	if ($x5 == x6$) go to PC+100	PC-relative branch if registers equal
	Branch if not equal	<code>bne x5, x6, 100</code>	if ($x5 != x6$) go to PC+100	PC-relative branch if registers not equal
	Branch if less than	<code>blt x5, x6, 100</code>	if ($x5 < x6$) go to PC+100	PC-relative branch if registers less
	Branch if greater or equal	<code>bge x5, x6, 100</code>	if ($x5 \geq x6$) go to PC+100	PC-relative branch if registers greater or equal
	Branch if less, unsigned	<code>bltu x5, x6, 100</code>	if ($x5 < x6$) go to PC+100	PC-relative branch if registers less
	Branch if greater/eq, unsigned	<code>bgeu x5, x6, 100</code>	if ($x5 \geq x6$) go to PC+100	PC-relative branch if registers greater or equal
Unconditional branch	Jump and link	<code>jal x1, 100</code>	$x1 = PC+4$; go to PC+100	PC-relative procedure call
	Jump and link register	<code>jalr x1, 100(x5)</code>	$x1 = PC+4$; go to $x5+100$	Procedure return; indirect call

RISC-V: 6 instruction formats

- R-Format: instructions using 3 register inputs
 - add, xor — arithmetic/logical ops
- I-Format: instructions with immediates, loads
 - addi, lw
- S-Format: store instructions: sw, sb
- SB-Format: branch instructions: beq, bge
- U-Format: instructions with upper immediates
 - lui — upper immediate is 20-bits
- UJ-Format: the jump instruction: jal

Name (Field Size)	7 bits	5 bits	5 bits	3 bits	5 bits	7 bits	Comments
R-type	funct7	rs2	rs1	funct3	rd	opcode	Arithmetic instruction format
I-type	immediate[11:0]		rs1	funct3	rd	opcode	Loads & immediate arithmetic
S-type	immed[11:5]	rs2	rs1	funct3	immed[4:0]	opcode	Stores
SB-type	immed[12,10:5]	rs2	rs1	funct3	immed[4:1,11]	opcode	Conditional branch format
UJ-type	immediate[20,10:1,11,19:12]				rd	opcode	Unconditional jump format
U-type	immediate[31:12]				rd	opcode	Upper immediate format ⁴⁰

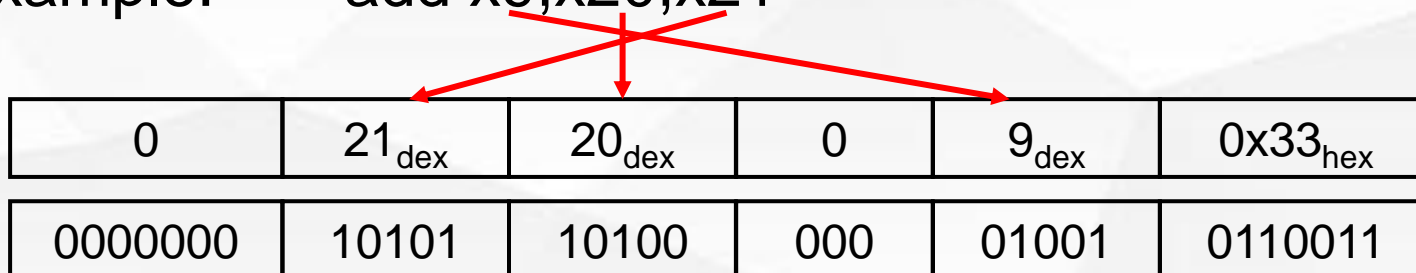
RISC-V R-format Instructions



• Instruction fields

- opcode: operation code
- rd: destination register number
- funct3: 3-bit function code (additional opcode)
- rs1: the first source register number
- rs2: the second source register number
- funct7: 7-bit function code (additional opcode)

• Example: add x9,x20,x21




0000 0001 0101 1010 0000 0100 1011 0011_{two} = 015A04B3₁₆

RISC-V I-format Instructions



- Immediate arithmetic and load instructions
 - rs1: source or base address register number
 - immediate: constant operand, or offset added to base address
 - 2s-complement, sign extended

• Example `addi x15, x1, -40`



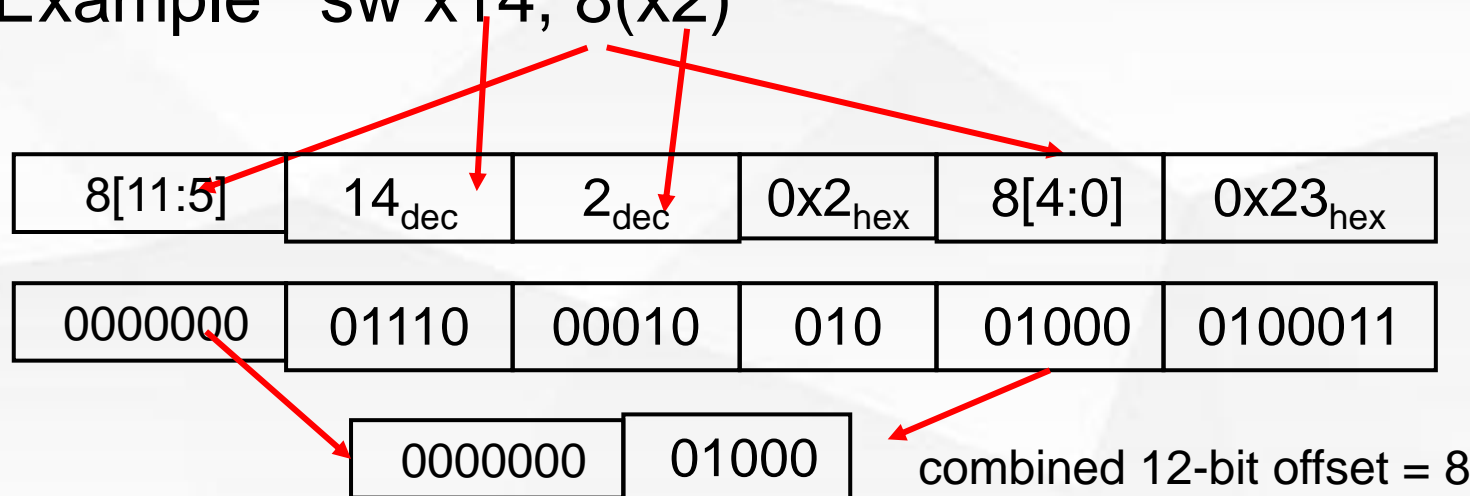
0xFD808793

RISC-V S-format Instructions



- Different immediate format for store instructions
 - rs1: base address register number
 - rs2: source operand register number
 - immediate: offset added to base address
 - Split so that rs1 and rs2 fields always in the same place

• Example `sw x14, 8(x2)`



immediate	rs1	funct3	rd	op
-----------	-----	--------	----	----

funct7	rs2	rs1	funct3	rd	op
--------	-----	-----	--------	----	----

immediate	rs1	funct3	rd	op
-----------	-----	--------	----	----

Register

Diagram illustrating the R-format instruction format. The instruction is a 32-bit word divided into six 5-bit fields: imm (5 bits), rs2 (5 bits), rs1 (5 bits), funct3 (3 bits), imm (5 bits), and op (5 bits). The PC (Program Counter) is shown as a 32-bit register that provides the rs1 and rs2 register indices to the ALU.

MIPS & RISC-V Instruction Formats

Register-register

	31	25	24	20	19	15	14	12	11	7	6	0															
RISC-V	funct7(7)					rs2(5)				rs1(5)				funct3(3)			rd(5)				opcode(7)						
	31	26	25	21	20	16	15	11	10	6	5	0															
MIPS	Op(6)					Rs1(5)					Rs2(5)					Rd(5)					Const(5)					Opx(6)	

Load

	31			20	19			15	14		12	11			7	6			0				
RISC-V	immediate(12)												rs1(5)			funct3(3)		rd(5)			opcode(7)		
	31			26	25			21	20					16	15				0				
MIPS	Op(6)				Rs1(5)				Rs2(5)				Const(16)										

Store

	31	25	24	20	19	15	14	12	11	7	6	0									
RISC-V	immediate(7)					rs2(5)				rs1(5)				funct3(3)		immediate(5)			opcode(7)		
	31	26	25	21	20	16	15						0								
MIPS	Op(6)					Rs1(5)				Rs2(5)				Const(16)							

Branch

	31	25	24	20	19	15	14	12	11	7	6	0									
RISC-V	immediate(7)					rs2(5)				rs1(5)				funct3(3)		immediate(5)			opcode(7)		
	31	26	25	21	20	16	15						0								
MIPS	Op(6)					Rs1(5)				Opx/Rs2(5)				Const(16)							

Concluding Remarks

- Design principles
 1. Simplicity favors regularity
 2. Smaller is faster
 3. Make the common case fast
 4. Good design demands good compromises
- Layers of software/hardware
 - Compiler, assembler, hardware
- MIPS: typical of RISC ISAs
 - c.f. x86