

# Lecture 10

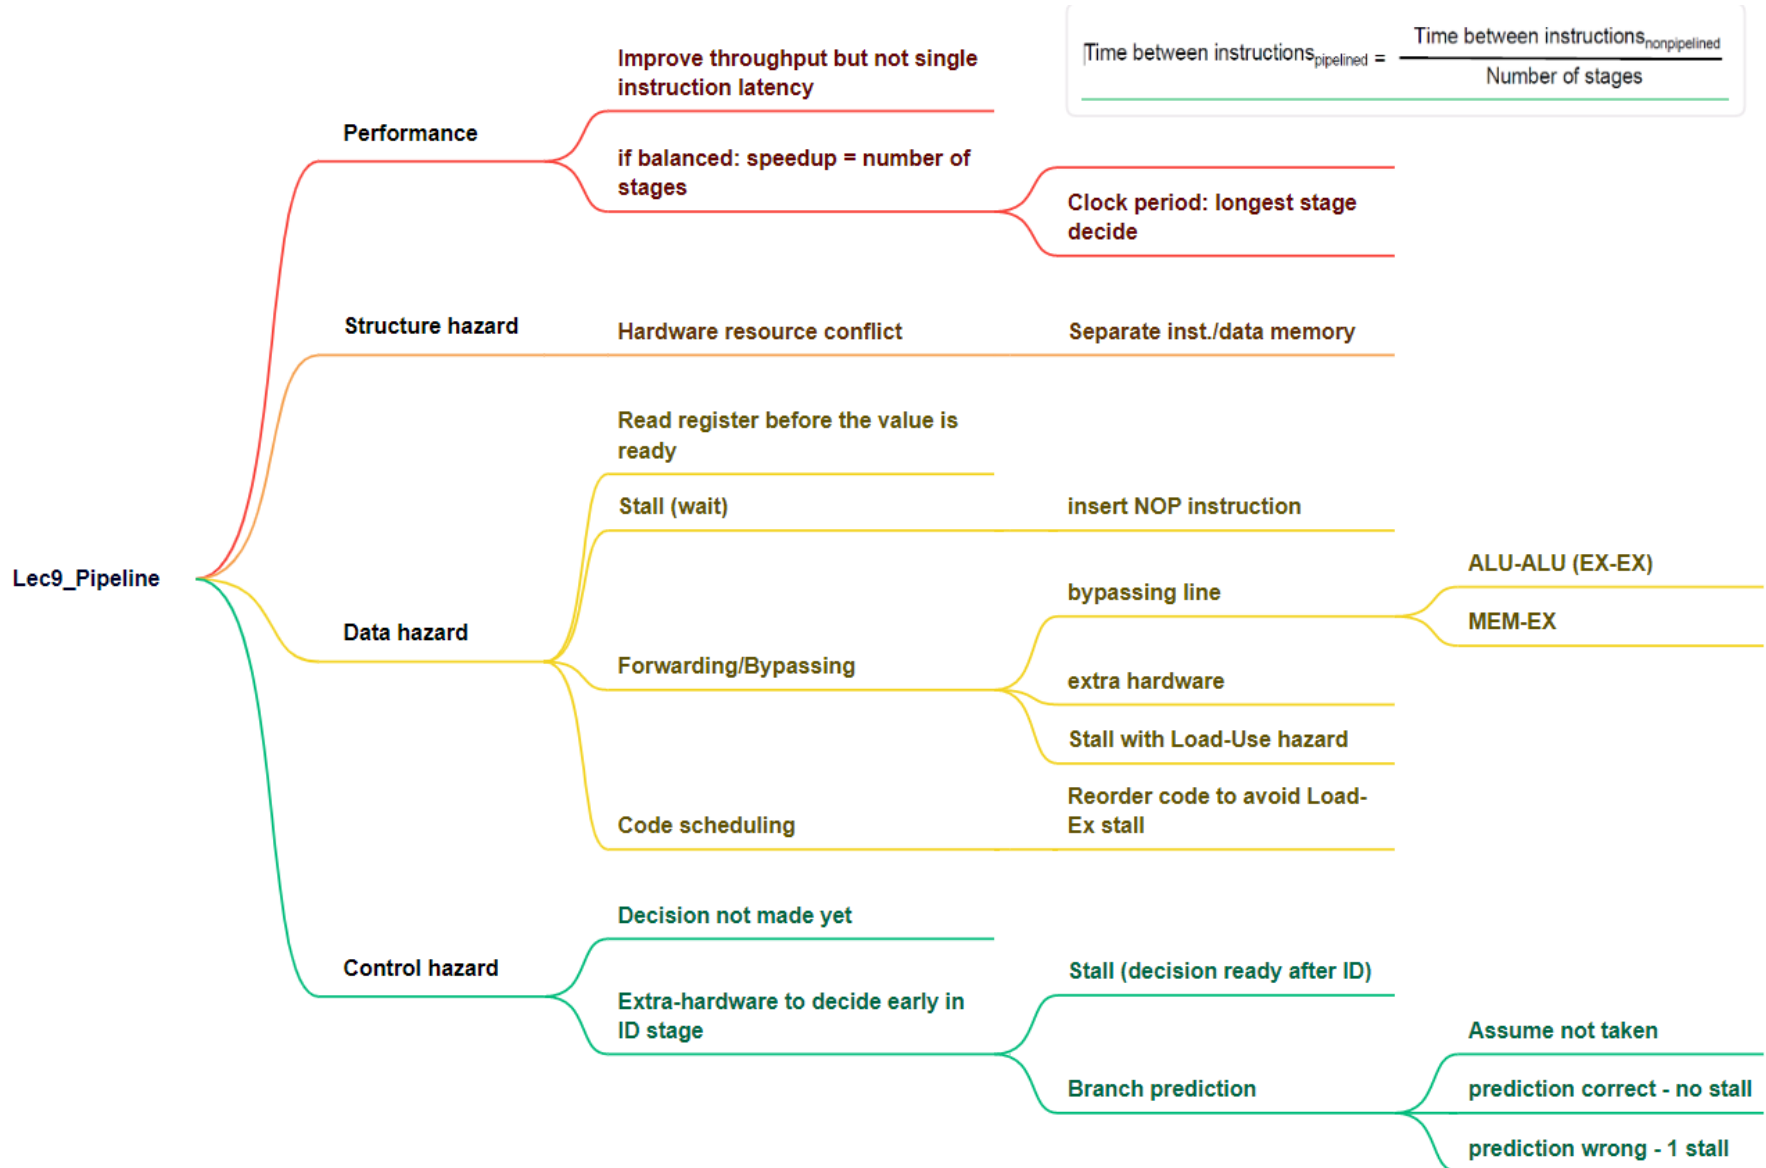
## Instruction-Level Parallelism

CS202 2023 Spring

# Today's Agenda

- Recap
  - Pipeline Performance
  - Pipeline Hazard
- Context
  - Instruction-level parallelism
  - Multiple issues
  - Scheduling
- Reading: Textbook 4.10-4.11, 4.14
- Self-study: The rest sections of Chapter 3


# Recap



$$\text{Time between instructions}_{\text{pipelined}} = \frac{\text{Time between instructions}_{\text{nonpipelined}}}{\text{Number of stages}}$$

# Instruction-Level Parallelism (ILP)

- Instruction-level parallelism: parallelism among instructions
  - Pipelining is one type of ILP: because pipeline executes multiple
- To increase ILP
  - Deeper pipeline
    - Less work per stage  $\Rightarrow$  shorter clock cycle
  - Multiple issue
    - Replicate pipeline stages  $\Rightarrow$  multiple pipelines
    - Start multiple instructions per clock cycle
    - $CPI < 1$ , so use Instructions Per Cycle (IPC)
    - E.g., 4GHz 4-way multiple-issue
      - 16 BIPS, peak  $CPI = 0.25$ , peak  $IPC = 4$
    - But dependencies reduce this in practice





# Multiple Issue

编译器把可以同时取到的指令，知道他们是否可以一起打包放到issue slots, 此过程是由编译器决定的，可以避免hazards  
在执行之前，就已经决定了

- **Static multiple issue (Compilers decide)**
  - Compiler groups instructions to be issued together
  - Packages them into “issue slots”
  - Compiler detects and avoids hazards
- **Dynamic multiple issue (Hardware decide)**
  - CPU examines instruction stream and chooses instructions to issue each cycle
  - Compiler can help by reordering instructions
  - CPU resolves hazards using advanced techniques at runtime

cpu决定的，不按照顺序执行，能执行那个执行那个，乱序执行，但是执行完后不会立刻放进寄存器

# Static Multiple Issue

- Compiler groups instructions into “issue packets”
  - Group of instructions that can be issued on a single cycle
  - Determined by pipeline resources required
- Think of an issue packet as a very long instruction
  - Specifies multiple concurrent operations
  - $\Rightarrow$  Very Long Instruction Word (VLIW)

相当于在一个clock里面，放置一个64bit的指令，由两个32bit的指令组成

# Scheduling Static Multiple Issue

- Compiler must remove some/all hazards
  - Reorder instructions into issue packets
  - No dependencies with a packet
  - Possibly some dependencies between packets
    - Varies between ISAs; compiler must know!
  - Pad with nop if necessary

packet之间不能有依赖关系  
同时执行的指令不能有依赖关系

# MIPS with Static Dual Issue

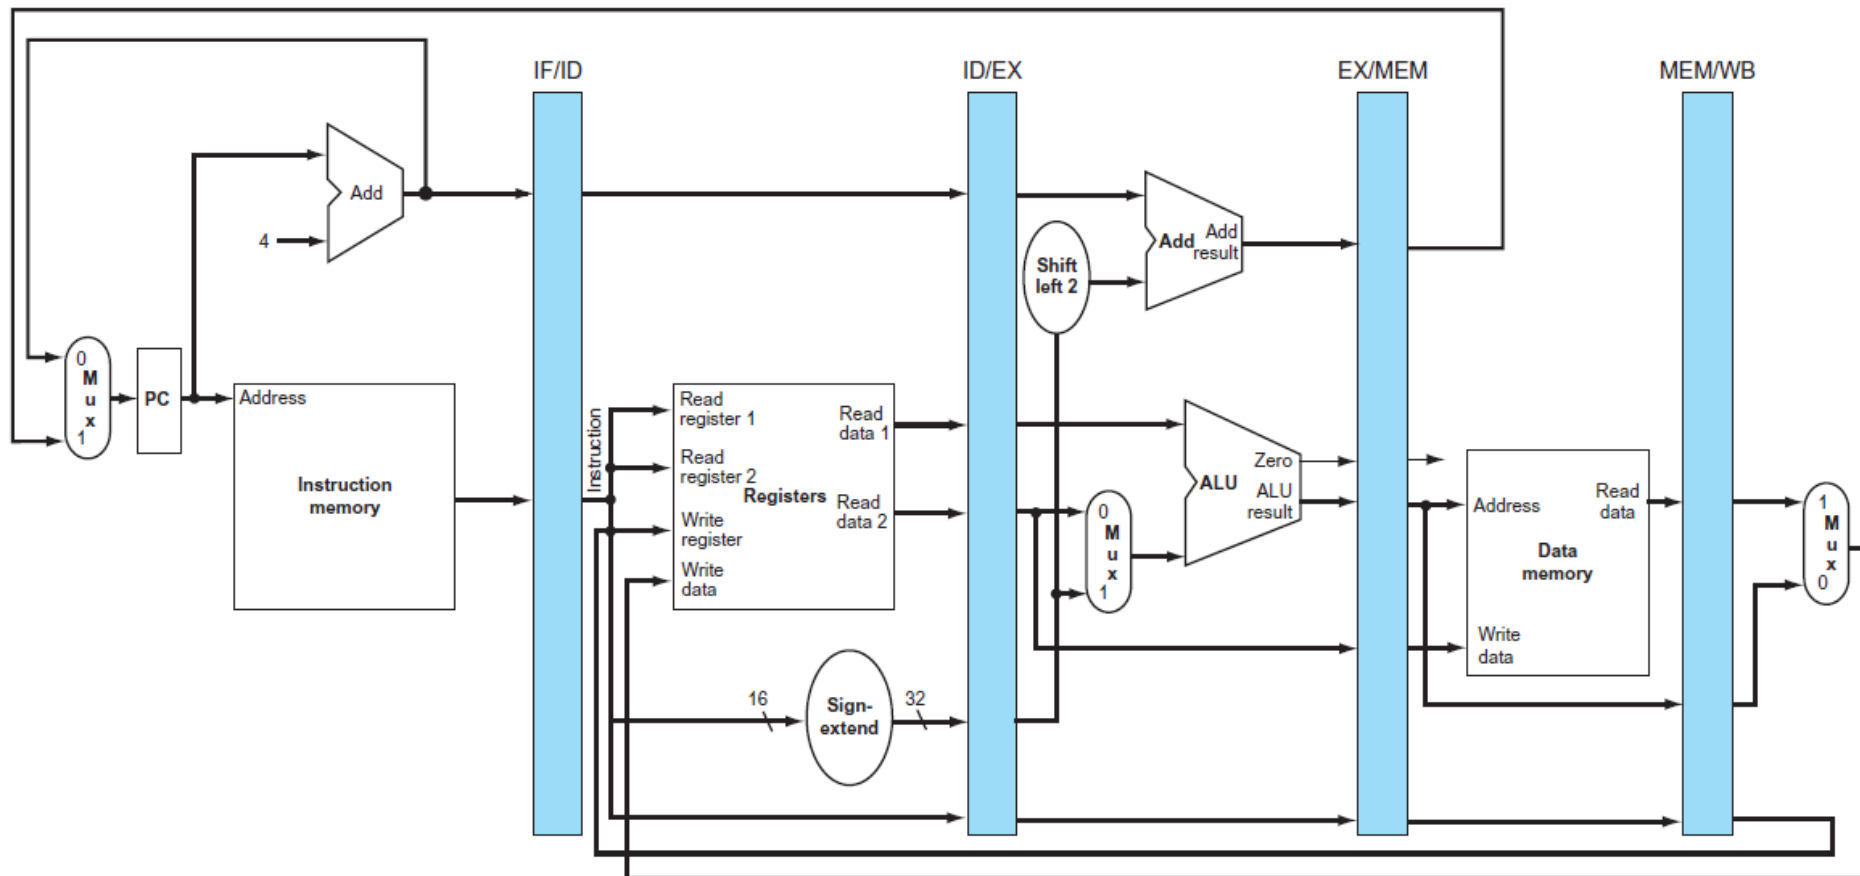
- Two-issue packets 一个packet里面，只能是alu/branch+la/sw 也就是每类各一个
  - Divide instructions into two types.
    - Type1: ALU/branch instruction
    - Type2: load/store instruction
  - During each cycle, execute a type1/type2 inst. simultaneously
  - 64-bit aligned instructions
    - Pad an unused instruction with nop

Address	Instruction type	Pipeline Stages						
		IF	ID	EX	MEM	WB		
n	ALU/branch	IF	ID	EX	MEM	WB		
n + 4	Load/store	IF	ID	EX	MEM	WB		
n + 8	ALU/branch		IF	ID	EX	MEM	WB	
n + 12	Load/store		IF	ID	EX	MEM	WB	
n + 16	ALU/branch			IF	ID	EX	MEM	WB
n + 20	Load/store			IF	ID	EX	MEM	WB



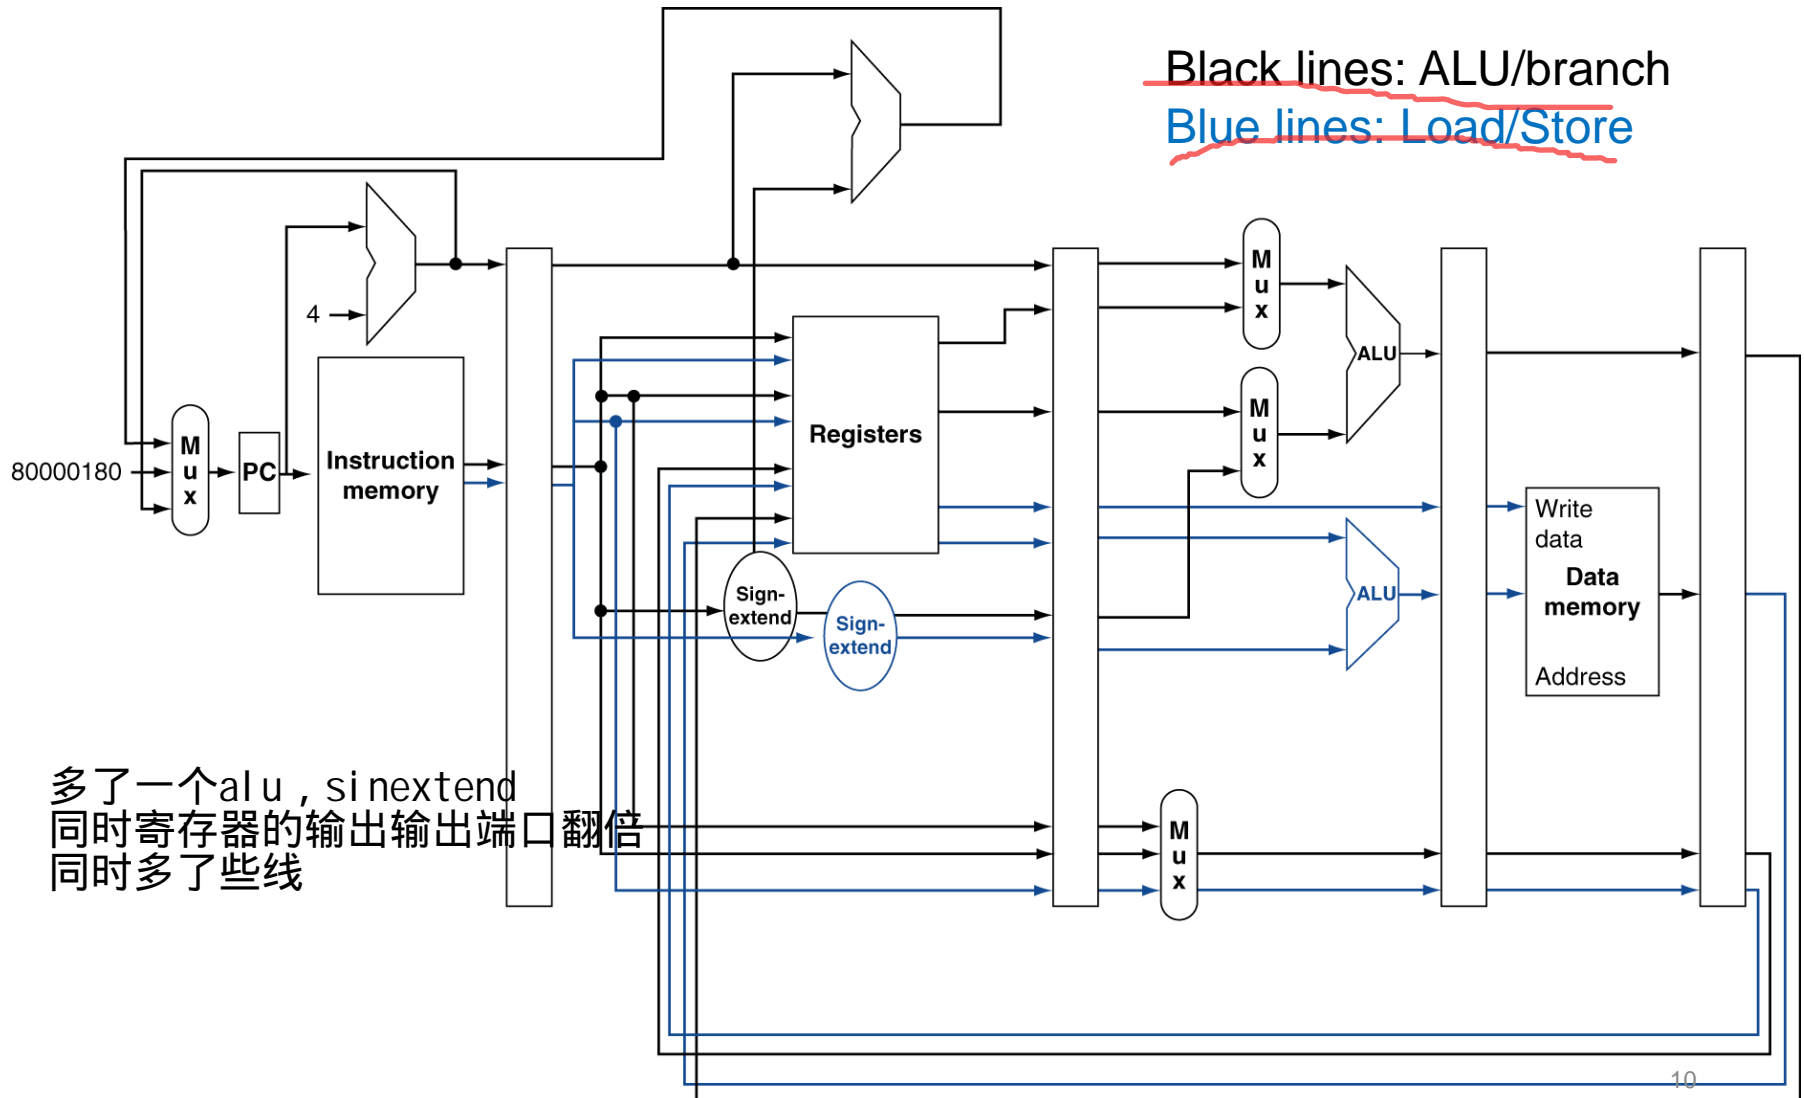
# A Basic MIPS Pipeline Datapath

把要做的事先缓存进来



# MIPS with Static Dual Issue

硬件上的改动



# Hazards in the Dual-Issue MIPS

- More instructions executing in parallel
- Removing hazards in static multiple issues
  - code scheduling and no-op insertion.
  - detect data hazard and generate all stalls between two issue packets, compiler avoids all dependencies within an instruction pair.
- EX data hazard
  - Forwarding avoided stalls with single-issue
  - Now can't use ALU result in load/store in same packet
    - add `$t0, $s0, $s1`      load `$s2, 0($t0)`
  - Split into two packets, effectively a stall
- Load-use hazard
  - Still one cycle use latency(number of clock cycles between load and use), but now two instructions
- More aggressive scheduling required

# Scheduling Example

- Schedule this for dual-issue MIPS

```
Loop: lw    $t0, 0($s1)      # $t0=array element
      addu  $t0, $t0, $s2    # add scalar in $s2
      sw    $t0, 0($s1)      # store result
      addi  $s1, $s1, -4     # decrement pointer
      bne   $s1, $zero, Loop # branch $s1!=0
```

	ALU/branch	Load/store	cycle
Loop:		lw    \$t0, 0(\$s1)	1
			2
			3
			4

# Scheduling Example

- Schedule this for dual-issue MIPS

```

Loop: lw    $t0, 0($s1)      # $t0=array element
      addu  $t0, $t0, $s2    # add scalar in $s2
      sw    $t0, 0($s1)      # store result
      addi  $s1, $s1, -4     # decrement pointer
      bne   $s1, $zero, Loop # branch $s1!=0
  
```

	ALU/branch	Load/store	cycle
Loop:		lw    \$t0, 0(\$s1)	1
			2
	addu  \$t0, \$t0, \$s2		3
			4

# Scheduling Example

- Schedule this for dual-issue MIPS

```

Loop: lw    $t0, 0($s1)      # $t0=array element
      addu  $t0, $t0, $s2    # add scalar in $s2
      sw    $t0, 0($s1)      # store result
      addi  $s1, $s1, -4     # decrement pointer
      bne   $s1, $zero, Loop # branch $s1!=0
  
```

	ALU/branch	Load/store	cycle
Loop:		lw    \$t0, 0(\$s1)	1
			2
	addu  \$t0, \$t0, \$s2		3
		sw    \$t0, 0(\$s1)	4

# Scheduling Example

## • Schedule this for dual-issue MIPS

```

Loop: lw    $t0, 0($s1)      # $t0=array element
      addu  $t0, $t0, $s2    # add scalar in $s2
      sw    $t0, 0($s1)      # store result
      addi  $s1, $s1, -4     # decrement pointer
      bne   $s1, $zero, Loop # branch $s1!=0
  
```

	ALU/branch	Load/store	cycle
Loop:		lw    \$t0, 0(\$s1)	1
	addi  \$s1, \$s1, -4		2
	addu  \$t0, \$t0, \$s2		3
		sw    \$t0, 4(\$s1)	4

bring addi to cycle 2 to avoid stall, but  
need to adjust sw offset

稍微改动下指令，先-4，再给他偏移回去

# Scheduling Example

- Schedule this for dual-issue MIPS

```

Loop: lw    $t0, 0($s1)      # $t0=array element
      addu  $t0, $t0, $s2    # add scalar in $s2
      sw    $t0, 0($s1)      # store result
      addi  $s1, $s1, -4     # decrement pointer
      bne   $s1, $zero, Loop # branch $s1!=0
  
```

	ALU/branch	Load/store	cycle
Loop:		lw    \$t0, 0(\$s1)	1
	addi  \$s1, \$s1, -4		2
	addu  \$t0, \$t0, \$s2		3
	bne   \$s1, \$zero, Loop	sw    \$t0, 4(\$s1)	4



# Scheduling Example

- Schedule this for dual-issue MIPS

```

Loop: lw    $t0, 0($s1)      # $t0=array element
      addu  $t0, $t0, $s2    # add scalar in $s2
      sw    $t0, 0($s1)      # store result
      addi  $s1, $s1, -4     # decrement pointer
      bne   $s1, $zero, Loop # branch $s1!=0
  
```

	ALU/branch	Load/store	cycle
Loop:	<b>nop</b>	lw    \$t0, 0(\$s1)	1
	addi  \$s1, \$s1, -4	<b>nop</b>	2
	addu  \$t0, \$t0, \$s2	<b>nop</b>	3
	bne   \$s1, \$zero, Loop	sw    \$t0, 4(\$s1)	4

# Scheduling Example

- Schedule this for dual-issue MIPS

```

Loop: lw    $t0, 0($s1)      # $t0=array element
      addu  $t0, $t0, $s2    # add scalar in $s2
      sw    $t0, 0($s1)      # store result
      addi  $s1, $s1, -4     # decrement pointer
      bne   $s1, $zero, Loop # branch $s1!=0
  
```

	ALU/branch	Load/store	cycle
Loop:	<code>nop</code>	<code>lw    \$t0, 0(\$s1)</code>	1
	<code>addi  \$s1, \$s1, -4</code>	<code>nop</code>	2
	<code>addu  \$t0, \$t0, \$s2</code>	<code>nop</code>	3
	<code>bne   \$s1, \$zero, Loop</code>	<code>sw    \$t0, 4(\$s1)</code>	4

$$\text{IPC} = 5/4 = 1.25 \text{ (c.f. peak IPC} = 2\text{)}$$

# Loop Unrolling

- Replicate loop body to expose more parallelism
  - Reduces loop-control overhead
- Use different registers per replication
  - Called “register renaming”
  - Avoid loop-carried “**anti-dependencies**”
    - Store followed by a load of the same register
    - Aka “**name dependence**” (Reuse of a register name)

Many anti-dependencies

```
for (i=4; i > 0;i--)
  a[i]=a[i]+5;
```

```
a[4]=a[4]+5;
a[3]=a[3]+5;
a[2]=a[2]+5;
a[1]=a[1]+5;
```

```
lw    $t0, 0($s1)
addu  $t0, $t0, $s2
sw    $t0, 0($s1)
addi  $s1, $s1, -4
```

a[4]=a[4]+5;

```
lw    $t0, 0($s1)
addu  $t0, $t0, $s2
sw    $t0, 0($s1)
addi  $s1, $s1, -4
```

a[3]=a[3]+5;

# Loop Unrolling Example

- Repeat the code in the loop

lw $\$t0$ , 0( $\$s1$ ) addu $\$t0$ , $\$t0$ , $\$s2$ sw $\$t0$ , 0( $\$s1$ )	lw $\$t1$ , -4( $\$s1$ ) addu $\$t1$ , $\$t1$ , $\$s2$ sw $\$t1$ , -4( $\$s1$ )	lw $\$t2$ , -8( $\$s1$ ) addu $\$t2$ , $\$t2$ , $\$s2$ sw $\$t2$ , -8( $\$s1$ )	lw $\$t3$ , -12( $\$s1$ ) addu $\$t3$ , $\$t3$ , $\$s2$ sw $\$t3$ , -12( $\$s1$ )
---	---	---	---

	ALU/branch 增添的	Load/store	cycle
Loop:	addi $\$s1$ , $\$s1$ , -16	lw $\$t0$ , 0( $\$s1$ )	1
	nop	lw $\$t1$ , 12( $\$s1$ )	2
	addu $\$t0$ , $\$t0$ , $\$s2$	lw $\$t2$ , 8( $\$s1$ )	3
	addu $\$t1$ , $\$t1$ , $\$s2$	lw $\$t3$ , 4( $\$s1$ )	4
	addu $\$t2$ , $\$t2$ , $\$s2$	sw $\$t0$ , 16( $\$s1$ )	5
	addu $\$t3$ , $\$t4$ , $\$s2$	sw $\$t1$ , 12( $\$s1$ )	6
	nop	sw $\$t2$ , 8( $\$s1$ )	7
	bne $\$s1$ , $\$zero$ , Loop	sw $\$t3$ , 4( $\$s1$ )	8

因为先做了s1的减法

# Loop Unrolling Example

- $IPC = 14/8 = 1.75$ 
  - Closer to 2, but at cost of registers and code size

	ALU/branch	Load/store	cycle
Loop:	addi <b>\$s1</b> , \$s1, -16	lw <b>\$t0</b> , 0(\$s1)	1
	<b>nop</b>	lw <b>\$t1</b> , 12(\$s1)	2
	addu <b>\$t0</b> , <b>\$t0</b> , \$s2	lw <b>\$t2</b> , 8(\$s1)	3
	addu <b>\$t1</b> , <b>\$t1</b> , \$s2	lw <b>\$t3</b> , 4(\$s1)	4
	addu <b>\$t2</b> , <b>\$t2</b> , \$s2	sw <b>\$t0</b> , 16(\$s1)	5
	addu <b>\$t3</b> , <b>\$t4</b> , \$s2	sw <b>\$t1</b> , 12(\$s1)	6
	<b>nop</b>	sw <b>\$t2</b> , 8(\$s1)	7
	bne <b>\$s1</b> , \$zero, Loop	sw <b>\$t3</b> , 4(\$s1)	8

# Dynamic Multiple Issue

- “Superscalar” processors
  - An advanced pipelining techniques that enables the processor to execute more than one instruction per clock cycle **by selecting them during execution.**
- CPU decides whether to issue 0, 1, 2, ... each cycle
  - Avoiding structural and data hazards
- Avoids the need for compiler scheduling
  - Though it may still help
  - Code semantics ensured by the CPU

不需要compiler做任何事，只需要硬件资源

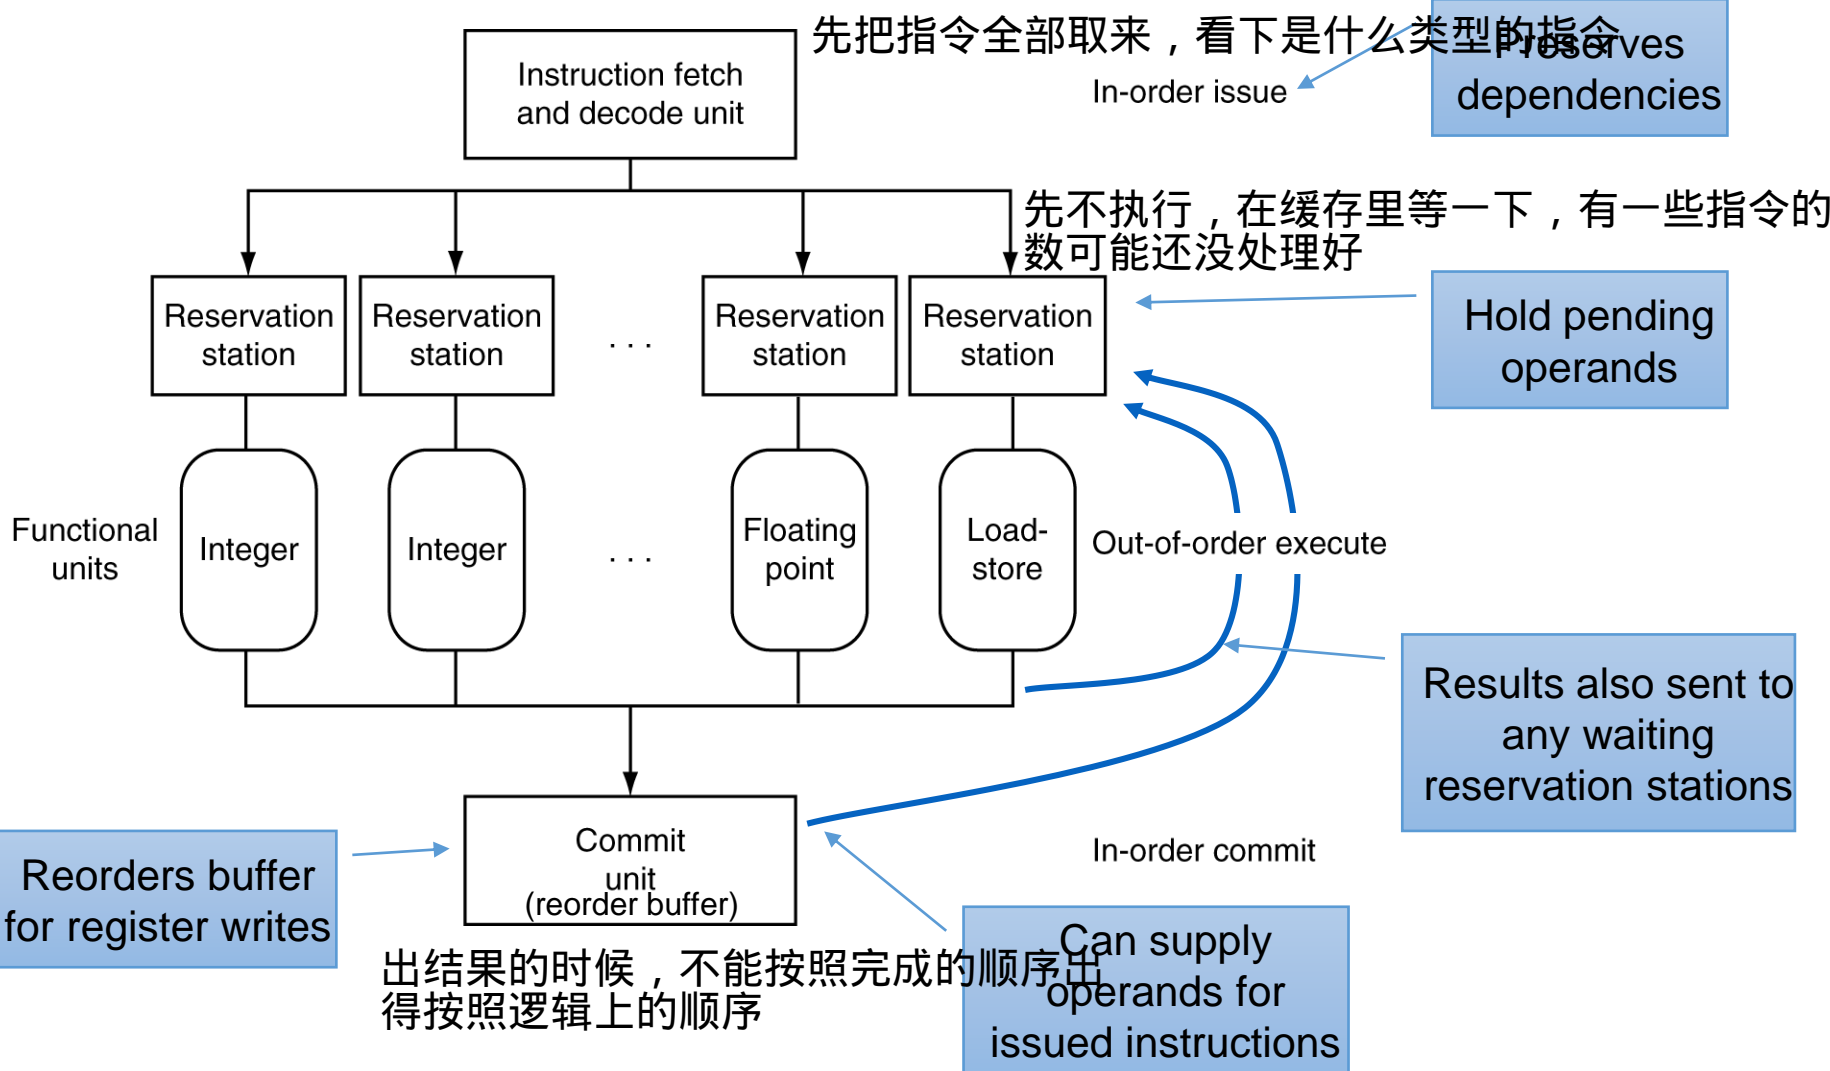
# Dynamic Pipeline Scheduling

- Hardware support for reordering the order of instruction execution
- Allow the CPU to **execute instructions out of order** to avoid stalls
  - But **commit result** to registers **in order**
- Example

```
lw      $t0, 20($s2)
addu    $t1, $t0, $t2
sub      $s4, $s4, $t3
slti    $t5, $s4, 20
```

  - Can start sub while addu is waiting for lw

# Dynamically Scheduled CPU





# Register Renaming

- Reservation stations and reorder buffer effectively provide register renaming
- On instruction issue to reservation station
  - If operand is available in register file or reorder buffer
    - Copied to reservation station
    - No longer required in the register; can be overwritten
  - If operand is not yet available
    - It will be provided to the reservation station by a function unit
    - Register update may not be required

# Why Do Dynamic Scheduling?

- Why not just let the compiler schedule code?
- Not all stalls are predictable
  - e.g., cache misses
- Can't always schedule around branches
  - Branch outcome is dynamically determined
- Different implementations of an ISA have different latencies and hazards

# Speculation for Static/Dynamic Multiple Issue

- “Guess” what to do with an instruction
  - Start operation as soon as possible
  - Check whether guess was right
    - If so, complete the operation
    - If not, roll-back and do the right thing
- Common to static and dynamic multiple issue
- Examples
  - Speculate on branch outcome
    - Roll back if path taken is different
  - Speculate on load
    - Avoid load and cache miss delay
    - Roll back if location is updated

# Compiler/Hardware Speculation

- Compiler can reorder instructions
  - e.g., move an instruction across a branch or a load across a store
  - Can include “fix-up” instructions to recover from incorrect guess
- Hardware can look ahead for instructions to execute
  - Buffer results and write to register or memory until it determines they are actually needed
  - Flush buffers on incorrect speculation

# Speculation and Exceptions

- What if exception occurs on a speculatively executed instruction?
  - e.g., speculative load before null-pointer check
- Static speculation
  - Can add ISA support for deferring exceptions
- Dynamic speculation
  - Can buffer exceptions until instruction completion (which may not occur)

# Summary

## • Techniques

- Static Multiple issue
- Dynamic Multiple issue
- VLIW
- Superscalar
- Loop unrolling
- Register renaming
- Dynamic scheduling
- Out-of-order execution
- Speculation

	<b>Static multiple issue</b>	<b>Dynamic multiple issue</b>
Decision made by	Compiler (software)	Processor (hardware)
Also called	Very long instruction word (VLIW)	Superscalar
Ways to remove hazard	Loop unrolling Register renaming Speculation	Dynamic scheduling Out-of-order execution Register renaming Speculation

# Does Multiple Issue Work?

- Yes, but not as much as we'd like
- Programs have real dependencies that limit ILP
- Some dependencies are hard to eliminate
  - e.g., pointer aliasing
- Some parallelism is hard to expose
  - Limited window size during instruction issue
- Memory delays and limited bandwidth
  - Hard to keep pipelines full
- Speculation can help if done well

# Power Efficiency

- Complexity of dynamic scheduling and speculations requires power
- Multiple simpler cores may be better



Microprocessor	Year	Clock Rate	Pipeline Stages	Issue width	Out-of-order/ Speculation	Cores	Power
i486	1989	25MHz	5	1	No	1	5W
Pentium	1993	66MHz	5	2	No	1	10W
Pentium Pro	1997	200MHz	10	3	Yes	1	29W
P4 Willamette	2001	2000MHz	22	3	Yes	1	75W
P4 Prescott	2004	3600MHz	31	3	Yes	1	103W
Core	2006	2930MHz	14	4	Yes	2	75W
UltraSparc III	2003	1950MHz	14	4	No	1	90W
UltraSparc T1	2005	1200MHz	6	1	No	8	70W




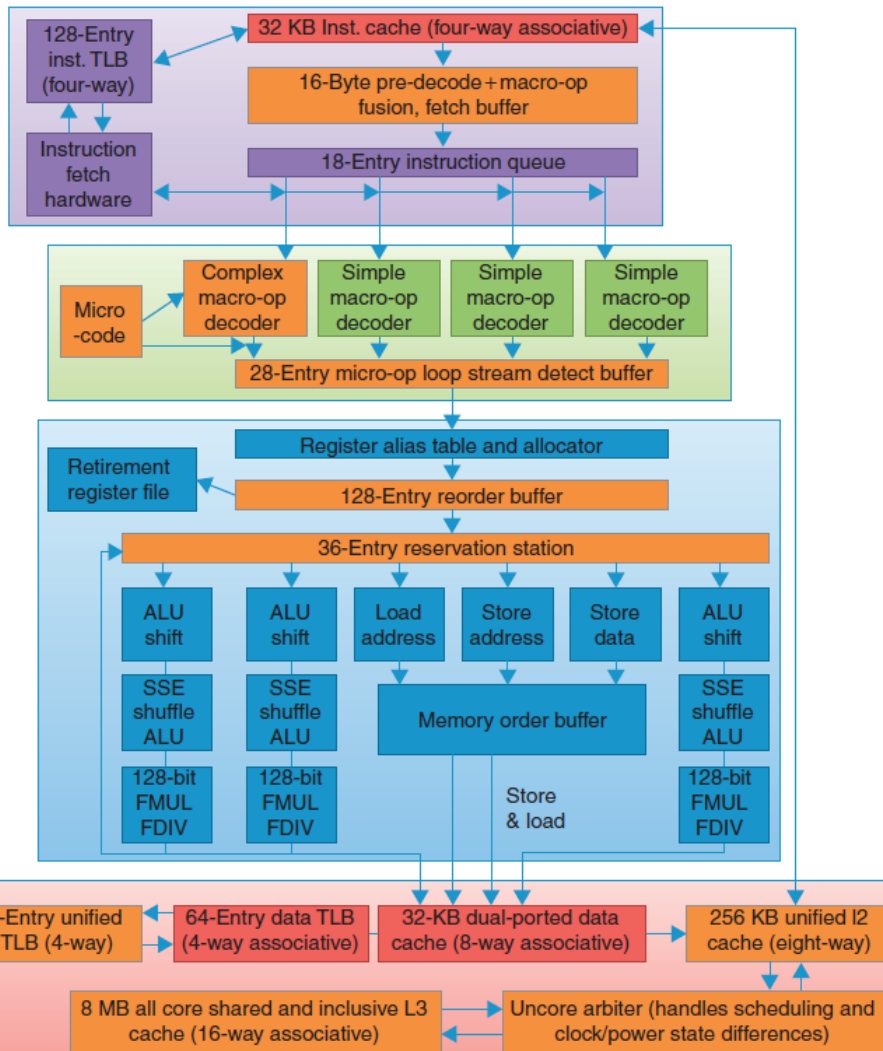
# Cortex A8 and Intel i7

Processor	ARM A8	Intel Core i7 920
Market	Personal Mobile Device	Server, cloud
Thermal design power	2 Watts	130 Watts
Clock rate	1 GHz	2.66 GHz
Cores/Chip	1	4
Floating point?	No	Yes
Multiple issue?	Dynamic	Dynamic
Peak instructions/clock cycle	2	4
Pipeline stages	14	14
Pipeline schedule	Static in-order	Dynamic out-of-order with speculation
Branch prediction	2-level	2-level
1 <sup>st</sup> level caches/core	32 KiB I, 32 KiB D	32 KiB I, 32 KiB D
2 <sup>nd</sup> level caches/core	128-1024 KiB	256 KiB
3 <sup>rd</sup> level caches (shared)	-	2- 8 MB

# ARM Cortex-A8 Pipeline & Performance



# Core i7 Pipeline & Performance



# Fallacies

- Pipelining is easy (!)
  - The basic idea is easy
  - The devil is in the details
    - e.g., detecting data hazards
- Pipelining is independent of technology
  - So why haven't we always done pipelining?
  - More transistors make more advanced techniques feasible
  - Pipeline-related ISA design needs to take account of technology trends

# Pitfalls

- Poor ISA design can make pipelining harder
  - e.g., complex instruction sets (VAX, IA-32)
    - Significant overhead to make pipelining work
    - IA-32 micro-op approach
  - e.g., complex addressing modes
    - Register update side effects, memory indirection
  - e.g., delayed branches
    - Advanced pipelines have long delay slots

# Concluding Remarks

- ISA influences design of datapath and control
- Datapath and control influence design of ISA
- Pipelining improves instruction throughput using parallelism
  - More instructions completed per second
  - Latency for each instruction not reduced
- Hazards: structural, data, control
- Multiple issue and dynamic scheduling (ILP)
  - Dependencies limit achievable parallelism
  - Complexity leads to the power wall